

HDFC Bank

Six digital journeys, four business objectives, one shared spine

35 use cases delivered into a live estate of 85 journeys and 150,000 users — by building the shared components once and letting each journey contribute only its boundary crossing.

6

journeys

35

use cases

4

business
objectives

5

shared
components

11


service-request
categories

2

legal entities
in one flow

BUSINESS OBJECTIVES

Every journey was commissioned against one of four business objectives, each with a named owner metric — no journey entered the backlog without one



The discipline this enforced

A journey with no owner metric is a feature request. Requiring the metric before the design meant scope arguments were settled by which option moved the number, not by who asked loudest — and it made the post-launch review a measurement exercise rather than a debate.



Six journeys and 35 use cases shipped into a live estate by building the shared spine once — the sixth journey cost roughly a quarter of the first

Situation

One of the largest digital journey estates in Indian banking: 85 journeys live across unassisted, assisted and partner channels, 100+ integrated applications, 150,000 platform users.

New journeys are never greenfield — each lands on customer, consent, document and status models already in production.

Complication

The six journeys crossed six different boundaries: payment rails, servicing systems, an external depository, a separate legal entity, customer state, and a regulatory obligation.

Built independently, each would re-implement identity, consent, documents and status — and each would need its own audit evidence and its own review slot.

Resolution

Treat the six as one delivery against four business objectives. Model the shared spine once; let each journey contribute only its boundary crossing.

Effort per journey fell from ~26 to ~4 person-weeks by journey six, and assurance review effort fell with it.

THE NUMBERS THAT MATTER

35

use cases across 6 journeys

6×

reduction in effort, journey 1 to 6

38%

of servicing demand deflected (target)

+27pp

funded-in-session rate (target)

11

request types moved to self-service

5

components reused by every journey

What this case demonstrates: designing for reuse inside a live estate where rework is the dominant cost · delivering across a legal-entity boundary under one consent model · governing a journey portfolio by owner metric rather than by backlog · separating programme context from personal contribution


Programme scale, personal contribution and modelled value are stated separately —
so every number can be checked against the right source

PUBLICLY REPORTED Programme context The environment the work ran inside. Retained with attribution, not restated as personal achievement. 150,000+ platform users 120M+ customers 9,000+ branches 85 journeys in the estate 100+ integrated applications 10,000+ daily reports generated	DELIVERY RECORD My ownership Work led or delivered directly during the 2024 engagement. Verifiable through the delivery record and references. 6 journeys, 35 use cases - Requirement through production - Journey and data-model design - Cross-entity flow with the insurer 11-category servicing taxonomy - Consent, audit and exception design	ILLUSTRATIVE Modelled value The business case structure and its arithmetic. Volumes are illustrative; the method and metric definitions are the substance. - Assumptions stated on slide 15 - Rates and multipliers, not claims - Baseline captured pre-release - Per-journey metric definitions - Replace with bank actuals - Method transfers unchanged	OUT OF SCOPE Not claimed Explicitly outside this case, stated so the boundary is unambiguous rather than left to inference. - The decade-long transformation - Enterprise lead-conversion movement - Enterprise NPS or cross-sell movement - Platform or architecture selection - Any confidential artefact or screen - Any figure without a named source
--	--	---	---

SITUATION

The estate already carried 85 journeys and 100+ integrated applications — so the binding constraint was fitting the existing spine, not building new capability

150K+	120M+	9,000+	100+	300+	25+	85	10,000+
platform users	customers	branches	applications integrated	integration touchpoints	applications decommissioned	journeys in production	daily personalised reports



Digital journeys live in the estate, by channel type


What a mature estate changes

- Identity, consent and document handling already run in production — re-implementing them is pure cost and a new audit surface.
- Status models are shared. A journey that invents its own states is invisible to enterprise reporting on day one.
- Every integration point is governed. New calls inherit an approval path rather than creating one.
- The scarce resource is not engineering capacity — it is change-window and assurance capacity.

The design question was never “how do we build this journey?” — it was “what here is genuinely new, and what already exists and must be reused?”

COMPLICATION

The leaks were not in the products — they were at the boundaries between the bank and everything on the other side of them



Illustrative relationship funnel, indexed to 100 accounts opened — the shape, not the values, is the point

The funding gap

An account opened and never funded is an acquisition cost with no relationship attached. The drop happens in the first session, not over weeks.

The servicing tax

Eleven request categories, each requiring an agent, each with a different owning system and a different definition of “done”.

The group boundary

The highest-value cross-sell sits in a sister entity — and the consent model, not the product, is what usually blocks it.

The compliance chase

A mandated declaration collected by manual follow-up consumes operations capacity indefinitely and still never reaches full coverage.

Each leak sits at a boundary the bank does not fully control — a payment rail, a back-office owner, an external registry, a separate regulated entity. That is why the six journeys were scoped by boundary rather than by product line.

Six journeys, six boundaries, four objectives — and one common requirement behind all of them

Journey	Boundary crossed	Objective	Use cases	Owner metric	Target
Digital account funding	Payment rails	GROW	5	Funded-in-session rate	+27pp
Loan servicing self-help	Servicing systems	SERVE	11	Deflection across 11 categories	38%
Depository KYC update	External registry	SERVE	4	Straight-through completion	72%
Group health insurance	Separate legal entity	CROSS	6	Offer-to-issued conversion	1.6×
Commercial card application	Customer-state boundary	GROW	5	Completion, ETB vs NTB	71 / 43%
Beneficial ownership declaration	Regulatory obligation	COMPLY	4	Self-declaration coverage	68%

The common requirement

Six different failure modes — but behind every one: prove who the customer is, capture consent, handle documents, track state, and survive an audit. That shared 80% is what justified building the spine before the second journey.


Sequencing rationale

The three SERVE and COMPLY journeys shipped first: lower approval load, faster to production, and they proved the shared spine in a live environment — which is what earned the assurance capacity the two revenue journeys needed.


Thirty-five use cases were delivered across the six journeys — and 28 of them reused the spine without adding a single new shared component



Digital account funding closes the gap between “account opened” and “account usable” — the point where new relationships silently stall



Step 05 is the hard one: the payment rail settles on its own timetable while the customer sits in a session that will expire.



% of newly opened accounts — calibrated to published abandonment and funding-uplift benchmarks

Reused from the spine: identity · consent · document store · exception queue · notifications · audit trail
insurance journeys

New to the spine: reversible-transaction state machine, inherited by the card and


The design problem

- Settlement finality and session lifetime are governed by different systems and different clocks.
- A pending state invisible to the customer produces a duplicate attempt and a support call.
- Reversal is not an edge case here — it is a first-class outcome that must be designed.


How it was resolved

- Funding modelled as its own state machine, not a field on the account record.
- Every intermediate state given a customer-visible label and expected resolution time.
- Reversal path built and tested before the happy path was signed off.
- Idempotency enforced at authorisation so a retry cannot double-debit.

Loan servicing self-help moved 11 request categories out of the agent queue — by fixing the taxonomy first and building the interface second



The work that made this succeed happened at step 01. Shipping the interface over an unfixed taxonomy would have moved confusion, not workload.



Deflection by category, calibrated to published benchmarks — the spread is the insight, not the average

38%


blended deflection across 11 categories

3


categories carry over half the total benefit

- ### Why the spread matters more than the average
- Document-retrieval requests deflect almost completely — no judgement, no exception, no downstream owner.
 - Money-movement requests (foreclosure, prepayment) barely deflect, and should not have been forced to.
 - A blended target would have hidden both facts and produced the wrong roadmap for phase two.
 - Eligibility was configurable per category, so a category could be opened or closed without a release.

Group health insurance crossed a legal-entity boundary — two regulators and two systems of record, presented to the customer as one continuous flow




Step 03 is where most cross-entity journeys fail: a single blanket consent is quicker to build and does not survive scrutiny from either regulator.



The design problem	How it was resolved	Why it mattered
- Two regulated entities, two supervisors, two definitions of who owns the customer record. - The customer experiences one journey; the data must not behave as though it is one entity. - A blanket consent is faster to build and fails the first time either supervisor examines it.	- Consent split: cross-entity data sharing separated from the insurance contract itself. - Field-level handover, purpose recorded per attribute — minimum necessary, not bulk. - System-of-record split: insurer owns the policy, bank owns the servicing view. - Underwriting is a black box by design — status shown, not reasoning.	- Group cross-sell is where bancassurance value sits — and where it usually leaks. - A defensible consent model is the difference between a launched journey and a withdrawn one. - The pattern generalises to any group-entity product sold through a bank channel.

Build effort fell roughly sixfold from the first journey to the sixth — the case for the seventh is now stronger than the case for the first ever was



Person-weeks per journey, illustrative model — spine investment front-loaded into journeys one and two

The argument this makes to a sponsor

The spine looks like overhead on journey one and pays for itself by journey three. Sponsors who fund journeys individually never see that curve, because each business case is written as though it were the first. Funding the portfolio rather than the journey is the decision that unlocks it.

26

person-weeks
journey 1

4

person-weeks
journey 6

6x

effort reduction
across the set

22


spine person-weeks
invested in total

What the curve does and does not prove

- It proves reuse held. A rising curve would have meant the spine was being forked under delivery pressure.
- It does not prove the journeys got easier — journey four was the hardest, and its bar reflects real new complexity.
- The spine cost is visible, not hidden: 22 person-weeks, paid mostly in the first two journeys.

THE OPERATING PATTERN

Five shared components carried all 35 use cases — each journey contributed only its boundary crossing, and none forked the spine



Why this is an operating-model decision, not a technical one

- Reuse survives only if someone owns the shared components and can refuse a journey-specific fork.
- That ownership must sit above any single journey's delivery date, or it loses every deadline argument.
- The governance cost is real and worth naming — it is paid in review time, not licence fees.

Where the payoff shows up

- Assurance effort per journey falls, because the control evidence is already established.
- Enterprise reporting stays intact — no journey introduces a private status vocabulary.
- The marginal journey gets cheaper, so the seventh business case is stronger than the first.

CONTROL

Six controls applied to every journey before release — designed in from the first sprint, because assurance capacity was scarcer than engineering capacity

	Account funding	Loan servicing	Depository KYC	Group health	Commercial card	Beneficial ownership
Consent Purpose-bound and timestamped	●	●	●	EXTENDED	●	●
Identity assurance Proportionate to the action, not the channel	EXTENDED	●	●	●	EXTENDED	●
Maker-checker Where an action is irreversible	EXTENDED	●	—	●	EXTENDED	—
Audit trail Continuous across the boundary	●	●	●	EXTENDED	●	●
Data minimisation Field-level, purpose recorded	●	●	●	EXTENDED	●	EXTENDED
Exception path A designed route back to a human	●	●	●	●	●	●

● standard control applied as-is EXTENDED control extended for this journey — not applicable

The cross-entity journey extended four of six controls — which is exactly where the design time went, and why it was the last of the six to ship despite being commissioned second. Designing controls in from sprint one was the speed strategy, not a constraint on it: a journey that arrives at review with controls bolted on goes to the back of the assurance queue.

The business case was built bottom-up per journey — volume × rate change × unit value — so every number traces to an assumption a sponsor can challenge

Journey	Volume assumption	Rate change modelled	Modelled monthly effect
Digital account funding	200,000 accounts / mo	34% → 61% funded in session	+54,000 funded / mo
Loan servicing self-help	420,000 requests / mo	38% blended deflection	≈17,300 agent-hours / mo
Depository KYC update	95,000 updates / mo	72% straight-through	≈68,000 branch touches removed
Group health insurance	85,000 offers / mo	2.1% → 3.4% issued	+1,100 policies / mo
Commercial card application	12,000 applications / mo	ETB 71% · NTB 43% completion	+1,900 completions / mo
Beneficial ownership	1.4M obligated customers	31% → 68% self-declared	≈518,000 chases avoided

How the model was used

- Each rate change was the journey's owner metric — the number the post-launch review measured.
- Volumes came from the bank; rate changes were the contested assumption and were negotiated explicitly.
- Sensitivity ran on the rate, never the volume — that is where disagreement sits.

Why it is shown as a model, not a result

Volumes and rates here are illustrative. What transfers is the structure: one owner metric per journey, a baseline captured before release, a rate change that is explicitly the thing being argued about, and an effect expressed in the unit the sponsoring function already manages — funded accounts, agent-hours, policies, chases. Substitute the bank's actuals and the arithmetic is unchanged.

Every journey was instrumented against a pre-release baseline — so the result could be attributed to the journey rather than to the quarter

Journey	Obj.	Primary metric	Baseline captured before release	Base	Target
Digital account funding	GROW	% funded within the opening session	Funded-within-7-days rate	34%	61%
Loan servicing self-help	SERVE	% of 11 categories completed without an agent	Agent-handled volume by category	0%	38%
Depository KYC update	SERVE	Straight-through completion rate	Branch-handled update volume	0%	72%
Group health insurance	CROSS	Offer-to-issued conversion	Assisted bancassurance conversion	2.1%	3.4%
Commercial card application	GROW	Completion by customer state	Assisted completion, split by state	52%	71 / 43%
Beneficial ownership	COMPLY	Self-declaration coverage	Manual chase completion rate	31%	68%

Why the baseline mattered more than the number

Six journeys launched into a bank running dozens of other changes in the same quarter. Without a per-journey baseline captured before release, any post-launch movement is arguable — and in a steering review an arguable number is worth less than no number at all. Instrumenting first is what makes the result defensible afterwards, and it is the discipline that transfers to any programme.

What held, what did not, and what I would do differently on the next estate of this size

What held

- Building the shared spine before journey two. By journey four the marginal cost of a new one had dropped roughly sixfold.
- Fixing the servicing taxonomy before building its interface — the single highest-return decision in the programme.
- Splitting consent across the entity boundary. It cost design time and saved far more in review.
- Designing the reversal path before the happy path on the funding journey.

What did not

- Cross-entity data agreements were run as a legal workstream alongside design. They are a design input and should have been sequenced first — it is why journey four shipped last.
- Eligibility rules were configurable but the approval path to change them was not, so configurability delivered less speed than promised.
- Assurance capacity was planned per journey rather than as a shared, scheduled constraint across all six.

What I would change

- Start with the counterparty, not the customer. The hardest constraints in five of six journeys came from the other side of the boundary.
- Treat the change-approval path as part of the configurability requirement — configurable but unapprovable is not configurable.
- Book assurance capacity as a programme-level resource sequenced with releases, not requested per journey.
- Publish the shared-spine contract early so journey teams design against it rather than negotiate with it.

The pattern behind all three columns

Everything that held was a sequencing decision taken early — spine before journey two, taxonomy before interface, reversal before happy path. Everything that did not hold was a dependency discovered late and then managed as a parallel workstream rather than as an input. On an estate this size the constraint is almost never capability; it is the order in which decisions are taken, and how early the hardest constraint is allowed to shape the design.

Four decisions carried the programme — and none of them were about the platform

- | | | |
|----|--|---|
| 01 | Govern the portfolio by owner metric | No journey enters the backlog without a named business objective and a single owner metric. Scope arguments are then settled by which option moves the number, not by who asks loudest. |
| 02 | Treat the boundary as the design unit | Journeys look unrelated on a roadmap and are structurally identical behind the boundary. Group them by what they cross and the reuse case makes itself — a sixfold effort reduction across six journeys. |
| 03 | Fix the taxonomy before the interface | A self-service front end over a broken request taxonomy relocates confusion to the customer and calls it deflection. The spread across the eleven categories, not the blended average, is what sets the next roadmap. |
| 04 | Make controls the speed strategy | In a large estate assurance capacity is scarcer than engineering capacity. Journeys arriving at review with controls designed in ship sooner than journeys that arrive fast and dirty. |

None of the four depend on the platform, the vendor or the bank. They are why the sixth journey cost roughly a sixth of the first — and why the seventh business case would have been the easiest one to write.