

Enterprise AI Platform RFP

Reading an enterprise AI RFP like an evaluator, not a bidder

A worked deconstruction of a multi-year, SI-scale enterprise AI tender at India's largest bank — reverse-engineering the scoring sheet before writing a word of response.

3 + 2

year term

~75

AI use cases

80:20

QCBS weighting

70%

technical qualifying bar

≈130

technical marks

14

organisations at
pre-bid

TENDER PARAMETERS · PUBLISHED IN THE RFP AND PRE-BID MINUTES

On-premise AI platform · ~75 GenAI and agentic use cases · 150–200 staff trained · managed services and capability transfer

Aditya Singh · Bid analysis authored for leadership bid / no-bid review · 2026

Most bid teams read an RFP as a list of requirements — an evaluator reads it as a scoring sheet with gates in front of it, and the two produce different bids

Situation

A landmark tender: a single system integrator to build an on-premise enterprise AI platform and deliver roughly 75 GenAI and agentic use cases over three years, extendable by two.

Six workstreams, five milestones, payment tied to each, and a competitive field spanning global SIs and BFSI-AI specialists.

Complication

Read as requirements, the document invites a capability response — everything the bidder can do. That response loses.

The marks are unevenly distributed, several are pass/fail gates that sit before the technical evaluation, and two published requirements contradict each other. None of that is visible from a requirements reading.

The method

Four passes in a fixed order: score the gates, map the marks, find the contradictions, then price. Only after all four is a bid / no-bid position meaningful.

Applied here, the method surfaced the go/no-go pivot on day one rather than three weeks in.

THE NUMBERS THAT MATTER

4

reading passes, in order

6

eligibility gates

≈130

technical marks mapped

3

published contradictions

40

marks in one cluster

1

genuine go/no-go pivot

The discipline this enforces: effort is committed only after the gates are cleared and the marks are mapped. A bid team that writes first and checks eligibility later spends its most expensive weeks on a response it may not be permitted to submit.

Built entirely from published tender documents — no confidential bid material, no commercial position, no client-identifying detail

TENDER DOCUMENTS

Published context

Facts published in the RFP, its appendices and the pre-bid minutes. Available to every bidder in the field.

- Scope, term and milestone structure
- Eligibility criteria and thresholds
- Technical scoring parameters
- Commercial structure and instruments
- Pre-bid attendance list
- Published timeline

BID ANALYSIS

My contribution

The four-pass method, the mark map, the contradiction test and the decision framework. Authored for leadership bid / no-bid review.

- Four-pass reading method
- Mark distribution reconstruction
- Gate assessment framework
- Contradiction identification
- Pre-bid question set
- Bid / no-bid decision structure

ILLUSTRATIVE

Modelled

Self-assessment scoring shown to demonstrate the method. Values are a worked example, not a real organisation's position.

- Self-scoring worked example
- Gate-clearance illustration
- Effort allocation by mark yield
- Sensitivity on qualifying threshold
- Replace with a real assessment
- Method transfers unchanged


OUT OF SCOPE

Not claimed

Deliberately excluded. This deck teaches a reading method; it is not a bid document and takes no commercial position.

- Any employer's actual bid position
- Any pricing or commercial strategy
- Whether a bid was submitted
- Any confidential internal assessment
- Any competitor's position
- The tender outcome

Six workstreams, five milestones and two price components — the shape of the deal before any judgement about whether to chase it




Payment released against milestone, with a share of each tied to SLA satisfaction

What the shape already tells you

- Sixty per cent of value sits behind use-case delivery and managed services — this is a run engagement wearing a build engagement's clothes.
- Only twenty per cent is released for standing the platform up, so platform cost must be recovered across three years of delivery.
- Independent QA at one-to-three and a standing red team are cost lines most capability responses forget to price.
- Intellectual property, the orchestration layer and all licences vest in the bank — which removes the usual product-annuity argument entirely.

Two clusters carry 40 of roughly 130 technical marks — and neither is the capability narrative most bidders lead with



Reconstructed distribution of technical marks from the published evaluation parameters

40

marks in architecture plus the presentation

50

marks awarded on counted track record

What the distribution actually implies

- Architecture and the live presentation are the same argument scored twice — write them together, not in separate workstreams.
- Fifty marks are counted, not argued: numbers of deployments, banks, GPU-stack builds, years of experience. No amount of writing moves them.
- That split decides effort allocation. Roughly half the score is fixed on the day the bid opens; the other half is genuinely winnable.
- Risk, governance and the site visit carry 15 marks and are usually assigned to whoever is least busy.

The allocation rule this produces

Spend proportionally to winnable marks, not to page count. If half the technical score is fixed by countable history, the bid team's effort belongs almost entirely in architecture, the implementation plan, the showcase use cases and the presentation — and the counted parameters belong in a partnering decision made before the bid starts, not in the writing.

Six pass/fail gates sit in front of the technical evaluation — and two of them decide the entire field before a single mark is awarded

Gate	Threshold as published	Nature	Decisive?
Prior platform track record	Three or more BFSI enterprise AI platform builds, each above a stated crore threshold	Countable history	Decisive
Named GPU-stack build	At least one successful enterprise AI platform on the named accelerator stack	Countable history	Decisive
Turnover	A three-year average well into the hundreds of crores	Audited financials	Decisive
Profitability	Profit before tax positive in two of the last three years	Audited financials	Standard
Client references	Three or more, a majority domestic, for enterprise AI platform delivery	Countable history	Standard
Entity and compliance	Domestic entity, procurement-policy compliant, no debarment or unsigned SLA overhang	Documentary	Standard

Why the gates come first	What these particular gates signal	The three routes out
- They are pass/fail. A bidder failing any one is never scored, however strong the response. - They are the cheapest thing in the document to assess — an afternoon, not a workstream. - They are also the only part that cannot be improved by effort during the bid window. Source: Published RFP and pre-bid minutes for the enterprise AI platform tender · author's assessment	- The turnover bar and the three-build threshold together pitch this squarely at the large system-integrator tier. - The named-stack requirement narrows the field again, to organisations with a specific prior build. - A specialist clears the capability test and fails the arithmetic — which is a partnering question, not a writing one.	- Bid direct, if every gate clears on the organisation's own numbers. - Bid as part of a consortium or under an SI, where the gates are met by the lead entity. - Decline, and redirect the bid team to a pursuit where the gates are winnable.

Three published requirements contradict each other — finding them is worth more than another page of capability narrative

01

Vendor neutrality versus a counted stack

THE DOCUMENT SAYS

The document requires an architecture with no OEM lock-in, and states that equivalents to the named accelerator stack are acceptable.

AND ALSO SAYS

A scoring parameter awards marks for the number of implementations on that specific named stack.

A bidder cannot optimise for both. Raising it at pre-bid forces a clarification that either opens the field or confirms it is closed — and either answer changes the bid decision.

02

Volume versus the delivery window

THE DOCUMENT SAYS

Roughly seventy-five use cases are to be delivered, in waves, over three years.

AND ALSO SAYS

The published phase-three window is materially shorter than that volume implies at any credible per-use-case effort.

Either the envelope is elastic or the effort assumption is much lower than a bidder would price. Clarifying it is the difference between a priced bid and a loss-making one.

03

Reuse-first versus new-build ownership

THE DOCUMENT SAYS

Bidders must maximise reuse of the bank's existing lakehouse, infrastructure and platforms.

AND ALSO SAYS

Intellectual property for anything newly built, plus the orchestration layer and licences, vests entirely with the bank.

The boundary between reused and newly supplied components is undefined, and it determines both the scope of the IP transfer and what the SI can carry to its next client.

A contradiction found before the clarification deadline is free. The same contradiction found after award is a change request nobody has budgeted for.

Scoring your own bid against the published sheet before writing turns a capability debate into an arithmetic one

Scoring parameter	Available	Realistic	Movable?	What would move it
Architecture & integration	20	16	Yes	Depth on MLOps, observability and the control tower
Technical presentation	20	16	Yes	Rehearsal, and the same story as the architecture section
Implementation approach	15	12	Yes	Wave plan, named resourcing, credible governance model
Risk, governance & site visit	15	12	Yes	Responsible-AI and model-risk mapping done properly
Three showcase use cases	10	8	Yes	Working demonstrations with quantified impact
AI-ML org experience	10	6	No	Years elapsed — fixed at bid open
Platform deploys — references	10	5	No	Prior delivery count — fixed
Deploys in banks	10	5	No	Prior delivery count — fixed
GPU-stack implementations	10	2	Partly	A partner or consortium lead who has them
GenAI / agentic experience	10	7	Partly	Production references, not pilots

Source: Author's method — self-assessment values are an illustrative worked example

8

A two-component price against a top-tier field — the commercial structure tells you who the tender was written for

Component A — fixed, three years

- Platform implementation
- Use-case implementation across the full programme
- Future use-case development
- Priced in full and evaluated in full

Component B — blended day rate

Managed services, training and handholding, priced against an indicative man-day pool taken for evaluation only — actual effort per use case is agreed jointly after award.

What the instruments signal

- A substantial earnest-money deposit and a five per cent performance guarantee — both scale-filters in themselves.
- Single currency, no deviations permitted: the response must be fully compliant or it is not evaluated.
- An indicative man-day pool used for comparison means the commercial score is partly a bidding artefact, not a cost estimate.
- Commercial envelopes are opened only for bidders who clear both eligibility and the technical bar.

WHO ATTENDED THE PRE-BID — FOURTEEN ORGANISATIONS

9 Global and large system integrators

5 BFSI-AI and product specialists

What the field composition actually tells a bidder

Nine global SIs alongside five BFSI-AI specialists is not a mixed field — it is two different bids being written against one document. The SIs will compete on scale, references and the counted parameters. The specialists can only win on the parts that are argued rather than counted: architecture depth, responsible-AI governance and demonstrable production use cases. A specialist writing a scale-and-breadth response is competing on the axis it cannot win.

The method's real deliverable is a short, high-yield question set — filed before the clarification deadline, when answers are still free

01	Reconcile neutrality with the counted stack	Does an equivalent implementation on a different accelerator stack score under the named-stack parameter, or only implementations on that specific stack?	<i>Opens or closes the field. Highest-value question in the set.</i>
02	Confirm the delivery envelope	Is the published phase-three window the full envelope for the stated use-case volume, or does it govern only the first wave?	<i>Determines whether the programme can be priced without a loss.</i>
03	Define the reuse boundary	Where does the boundary sit between reused bank components and newly supplied ones, and how does that map to the IP transfer?	<i>Determines scope, IP obligations and what carries to the next client.</i>
04	Clarify infrastructure ownership	Is GPU infrastructure supplied by the bank, or must the bidder size, price and supply it?	<i>Materially changes the capital shape of the bid.</i>
05	Fix the managed-services term and exit	What is the term of managed services, and what are the exit criteria for capability transfer?	<i>Determines the annuity and the end-state obligation.</i>
Five questions, filed on time, are worth more to a bid than a hundred additional pages of capability narrative.			

Four rules for reading any large competitive tender

- | | | |
|----|---|---|
| 01 | Score the gates in the first afternoon | Pass/fail criteria are the cheapest part of the document to assess and the only part effort cannot change. Assess them before a bid team is stood up, not after three weeks of writing. |
| 02 | Reconstruct the scoring sheet before writing | Marks are unevenly distributed and a large share is usually fixed by countable history. Effort should follow winnable marks, not page count or requirement count. |
| 03 | Hunt the contradictions | Large documents are written by several people. Where two published requirements cannot both be optimised for, the clarification is free before the deadline and expensive after award. |
| 04 | Score yourself before you argue | A self-assessment against the published sheet turns a capability debate into arithmetic — and shows in week one whether the gap is a writing problem or a partnering decision. |

The method is four passes in a fixed order: gates, marks, contradictions, price. Reversing the order is how bid teams spend their most expensive weeks on a response they were never permitted to submit.