

Punjab National Bank

The RM's day, rebuilt — from a system of record to a system of action

A relationship manager does not need more information. They need the day sequenced, the admin removed, and the next action already decided — which is an operating-model change wearing a workbench.

3

constructs, one
workbench

18

capabilities
mapped

1

question the day
plan answers

4

channels, one
conversation

0

systems the RM
leaves

1

adoption test that
matters

DEMONSTRATION SCOPE · ALIGNED TO THE BANK'S BUSINESS VERTICALS

Relationship-manager cockpit · unified engagement layer · intelligence in the flow of work

Aditya Singh · Solution analysis authored for demonstration design and executive briefing · 2026

Relationship-manager platforms fail by adding screens — the ones that work remove admin, sequence the day and put the decision where the work happens

Situation

A large bank with a substantial relationship-manager population already has a customer platform. It holds the data. Relationship managers can find anything in it if they know where to look.

The prioritised capabilities were a cockpit, a unified engagement layer and an intelligence layer.

Complication

A system of record answers questions. A relationship manager's problem is not unanswered questions — it is that the day is unsequenced, the admin is heavy, and the context for a conversation has to be assembled manually before every call.

Adding capability to a system of record produces a richer system of record. Adoption stalls, because nothing about the day got easier.

The reframe

Design for the working day rather than for the data model. Three constructs: a cockpit that sequences the day and answers what should I do next; an engagement layer that makes every channel one conversation; and intelligence that arrives inside the task rather than in a separate screen.

The test for every capability is whether it removes a step, not whether it adds a view.

THE NUMBERS THAT MATTER

3

constructs

18

capabilities mapped

1

test applied to each

4+

channels, one thread

2

clocks feeding the plan

1

measure that predicts adoption

The judgement this argues: relationship-manager tooling is an operating-model intervention, not a software deployment. The capability list is the easy part — the hard part is that a relationship manager will only adopt something that makes today shorter, and most features make today longer before they make it better.

The capability scope is the client's; the operating-model argument, adoption analysis and measurement design are mine

DEMONSTRATION SCOPE

Client context

The three constructs and their capabilities as scoped by the bank and aligned to its business verticals.

- Cockpit, engagement and intelligence
- Eighteen named capabilities
- Existing customer platform in place
- Multi-channel engagement requirement
- Household and family linkage
- Reporting and productivity expectation

SOLUTION ANALYSIS

My contribution

The operating-model reframe, the capability test, the adoption analysis and the measurement design.

- Day-shaped design argument
- Removes-a-step capability test
- Three-construct mapping
- Adoption failure-mode analysis
- Measurement and leading indicators
- Demonstration narrative design

ILLUSTRATIVE

Modelled

Time-allocation and adoption figures illustrate the argument. Directional, not measured.

- Day time-allocation split
- Adoption curve by capability type
- Step-removal estimates
- Replace with bank actuals
- Shape is the argument
- Method transfers unchanged

OUT OF SCOPE


Not claimed

Excluded deliberately. A pre-demonstration solution analysis, not a delivery record.

- Delivery of this capability
- Any customer or portfolio data
- The bank's productivity figures
- Any commercial position
- Whether the pursuit proceeded
- Benchmark figures as commitments

THE PROBLEM

The constraint is not information — it is that selling time is the smallest slice of the day and the easiest to give away



What the shape implies for design

- Published benchmarks put commercial-banking RMs at 25–30% of time in client dialogue. Everything else is preparation, recording or coordination.
- Context assembly before a call is the least visible cost and the most automatable — it is the same lookups every time.
- Admin and logging is the largest slice, and the one relationship managers resent most, because it produces nothing they value.
- A capability that adds a screen competes for the twenty-four per cent. A capability that removes a step gives some of the rest back.

Time allocation calibrated to published commercial-banking benchmarks: 25–30% of an RM's time in client dialogue

The test applied to every capability in scope

Does this remove a step from the day, or does it add a view to the platform?

Removing a step: a call-preparation brief generated before the meeting, a disposition captured once and reused everywhere, a summary written automatically after a conversation. The relationship manager does less and the record improves.

Adding a view: another dashboard, another report, another place to look. Each is individually useful and collectively they are why relationship managers keep their own spreadsheets. Both kinds of capability appear in every requirement list, and separating them is the single most useful thing to do before designing a demonstration.

A cockpit that sequences, an engagement layer that unifies, and intelligence that arrives inside the task — each answers a different part of the day

Relationship cockpit

What should I do next?

- One-glance day plan across portfolio, tasks and alerts
- Prioritised engagements — maturities, renewals, service commitments
- Click-to-call and one-touch actions on the plan itself
- Live performance against branch and portfolio targets
- Daily productivity and pipeline reporting generated, not compiled

Unified engagement layer

What has this customer already been told?

- One conversation history across branch, phone, email and messaging
- Customer view with household and family linkage
- Live messaging for real-time service in the same thread
- Structured dispositions capturing every interaction once
- Interaction summaries generated automatically after contact

Intelligence in the flow

What is the right thing to say?

- Next-best-action ranked by value, arriving in the task not a report
- Predictive scoring for lead quality and attrition risk
- Call-preparation briefs assembled before every conversation
- Pre-qualified, eligibility-checked offers rather than product lists
- Guided playbooks and compliant scripts at the point of contact
- An assistant for policy-accurate answers during a call

Why the third construct only works because of the first two

Intelligence delivered into an unsequenced day is another notification. A next-best-action that arrives without the day being prioritised competes with everything else the relationship manager is already behind on, and loses.

Sequencing has to come first, engagement history second, and only then does a recommendation have somewhere to land. That ordering is the difference between an intelligence layer that changes behaviour and one that adds to the noise — and it is why the three were scoped as one programme rather than three.

Relationship managers judge a platform in the first week — and most capabilities make the day longer before they make it better

Capability	Effect in week one	Adoption outcome
Auto-generated call-preparation brief	Saves the lookups they were already doing manually	Immediate — used from day one
Automatic interaction summary	Removes the write-up they most resent	Immediate — strongest single driver
Click-to-call from the plan	Removes a copy, a paste and a dial	Immediate — small and constant
Prioritised day plan	Replaces their own list, which they trust more at first	Slow — needs a week of being right
Next-best-action recommendations	Adds a decision to evaluate, until it is trusted	Slow — earns trust or is ignored
Structured dispositions	Adds a field to complete on every interaction	Resisted — unless it removes reporting later

Sequence the rollout by week-one effect

- Lead with the three that save time immediately. They buy the credit needed for everything else.
- Introduce the day plan once there is enough interaction history for it to be right more often than their own list.
- Recommendations last, once the plan is trusted.

The one that always causes trouble

- Structured dispositions cost the relationship manager time and benefit the organisation.
- They are only accepted if something visible is removed in return — usually the daily or weekly report they currently write by hand.
- Shipping it without removing that report is how you lose the room.

What predicts adoption

- Not training hours and not feature count. The number of steps removed from the first hour of the day.
- Measure it directly: how long from logging in to the first customer contact.
- It moves before any commercial metric and is the earliest honest signal.

Four leading indicators that move before any commercial number — and two lagging ones that only become attributable once the leading ones have moved

Measure	Why it is the right one	Baseline needed	Type
Login to first customer contact	Directly measures whether the day starts faster	Current elapsed time, sampled	Leading
Share of plan actioned	Shows whether the sequencing is trusted, not just delivered	Not applicable — new	Leading
Interactions logged per contact	Rises when logging becomes automatic rather than manual	Current logging completeness	Leading
Context lookups before a call	Falls when preparation is generated rather than assembled	Observed count per call	Leading
Contacts per relationship manager per day	The productivity claim — but only credible once the four above move	Current daily contact rate	Lagging
Conversion on recommended actions	The intelligence claim — needs several months of loop closure	Conversion on self-selected actions	Lagging

Why the leading and lagging split matters when the pilot is reviewed

A relationship-manager platform is usually reviewed at three months against a commercial number, and at three months the commercial number has not moved — because behaviour change precedes revenue and the sample is small.

Agreeing before the pilot that the first review is judged on the four leading measures, and the commercial claim on the two lagging ones at a later point, is what prevents a working programme being cancelled at the first checkpoint. It also happens to be the honest way round: if the leading measures have not moved, the platform is not being used, and no amount of waiting will produce the lagging ones.

Four rules for frontline tooling that actually gets used

- | | | |
|----|--|---|
| 01 | Design for the day, not the data model | A system of record answers questions. A relationship manager's problem is an unsequenced day and heavy admin. Adding capability to a record system produces a richer record system and no behaviour change. |
| 02 | Apply the removes-a-step test | Every requirement list mixes capabilities that remove work with capabilities that add a view. They look alike on paper and behave completely differently in week one. Separate them before designing anything. |
| 03 | Sequence rollout by week-one effect | Lead with what saves time immediately, because that buys the credit needed for the things that take a week to prove. Recommendations land last, once the day plan is already trusted. |
| 04 | Agree the review basis before the pilot | Behaviour change precedes revenue. Judging a three-month pilot on a commercial number cancels working programmes; judging it on leading indicators is both fairer and more honest, because flat leading indicators mean it is not being used. |

The capability list was never the hard part. A relationship manager will adopt something that makes today shorter and ignore something that makes today richer — and almost every frontline platform is designed for the second.