Abu Dhabi Islamic Bank

Eight programmes, three years — what compounds when one team owns the sequence

A wholesale and business banking transformation delivered as eight sequential programmes rather than one. The interesting result is not any single programme; it is what the eighth cost compared with the first.

8	3	5	4	1	0
programmes delivered	years, one owning team	control functions coordinated	systems of record integrated	shared spine built once	programmes restarted

ENGAGEMENT SCOPE · WHOLESALE, BUSINESS AND FINANCIAL-INSTITUTION BANKING

Customer 360 and relationship management · deal and credit approval · service and complaints · campaigns · mobility · account opening · account planning · early warning

Aditya Singh · Project Lead and Product Manager · 2018–2020

Vendor and platform names removed. Programme scope is the client's; the sequencing analysis and the compounding argument are mine.

EXECUTIVE SUMMARY

Eight programmes over three years, sequenced so each one paid for part of the next — the order was the design decision, not the scope

Situation

A Gulf Islamic bank transforming wholesale, business and financial-institution banking. Eight distinct programmes, each with its own sponsor, business case and go-live.

One delivery team carried all eight, across three years, on one platform.

Complication

Eight sponsors means eight business cases, each written as though it were the only programme. Left to that, each would have built its own customer view, its own approval workflow, its own alerting and its own reporting.

That is the default outcome in a multi-sponsor portfolio, and it produces eight systems on one platform.

What was done instead

Sequence deliberately so the foundational programme lands first, and every subsequent one extends the shared spine rather than building beside it.

Customer 360 first, then everything that needs a customer view. Approvals second, then everything that needs a control decision. By programme six, most of what a new programme needed already existed.

THE NUMBERS THAT MATTER

8	**3**	**5**	**4**	**2**	**1**
programmes, one team	foundational, five extending	control functions in one flow	systems of record	reusable spines built	sequencing decision

Why the sequence mattered more than any single programme: in a multi-sponsor portfolio nobody owns the order. Each sponsor optimises their own programme, and the platform fragments quietly across three years. Owning the sequence is the only lever that prevents it — and it is available exactly once, at the start.

Programme scope and outcomes come from the delivery record; the sequencing argument and the effort curve are analysis

CLIENT SCOPE

Programme context

The eight programmes, their scope and the systems integrated, as defined and accepted by the client.

- Eight programmes, 2018–2020
- Wholesale, business and FI banking
- Five control functions in scope
- Four systems of record integrated
- Mobile and branch channels
- Group and related-party model

DELIVERY RECORD

My ownership

Project Lead and Product Manager across the portfolio — requirement through production on each programme.

- All eight, requirement to production
- Portfolio sequencing decisions
- Parallel-approval design
- Early-warning governance pattern
- Group and related-party data model
- Control-function engagement

ILLUSTRATIVE

Modelled

The effort curve and reuse percentages illustrate the compounding argument. Directional, not measured.

- Effort index per programme
- Reuse share by programme
- Sequencing counterfactual
- Replace with programme actuals
- Shape is the argument
- Method transfers unchanged

OUT OF SCOPE

Not claimed

Excluded deliberately. A delivery retrospective, not a business-outcome claim or a bank disclosure.

- Bank financial or commercial outcomes
- Any confidential client document
- Any screen, flow or requirement text
- Credit decisions or portfolio data
- Independently audited results
- Commercial terms of engagement

Source: Author's delivery record (ownership) · author's synthesis (sequencing argument and models)

Eight programmes in three groups — three built foundations, three built on them, and two extended the platform into governance

Foundations programmes 1–3

- Enterprise Customer 360 and relationship management across wholesale, business and financial-institution banking
- Wholesale deal and credit approval, with five control functions approving in parallel
- Enterprise service and complaint management, taxonomy-driven with queues and escalation

Built on the foundation programmes 4–6

- Marketing, prospecting and campaign management with segmentation and campaign lifecycle
- Mobile relationship-manager enablement across managed devices on both mobile platforms
- Account opening with automated control-function approvals — legal, compliance and Sharia

Governance extensions programmes 7–8

- Target setting and strategic account planning — planned against actual, wallet share, seasonality
- Early-warning integration and risk-alert governance, with human review and justification

Why Customer 360 had to be first

- Every subsequent programme needed a customer, a group hierarchy and an entitlement model.
- Built second, it would have had to reconcile with whatever the first programme invented.
- It is the only programme in the list with no useful business case of its own — which is exactly why it is usually sequenced last.

Why approvals had to be second

- Account opening, deal management and early warning all need a control-function decision recorded and evidenced.
- Building the parallel-approval engine once, in the programme that most needed it, meant three later programmes inherited it.
- The alternative is three approval workflows with three audit trails.

What that left for the rest

- Campaigns, mobility, planning and early warning each contributed their own domain logic and almost nothing else.
- Each still had a real business case and a real sponsor — the reuse was invisible to them.
- That invisibility is the point: sponsors fund outcomes, not spines.

Source: Author's delivery record, 2018–2020 engagement · author's analysis

By programme six most of what a new programme needed already existed — the reuse share rose faster than the scope fell

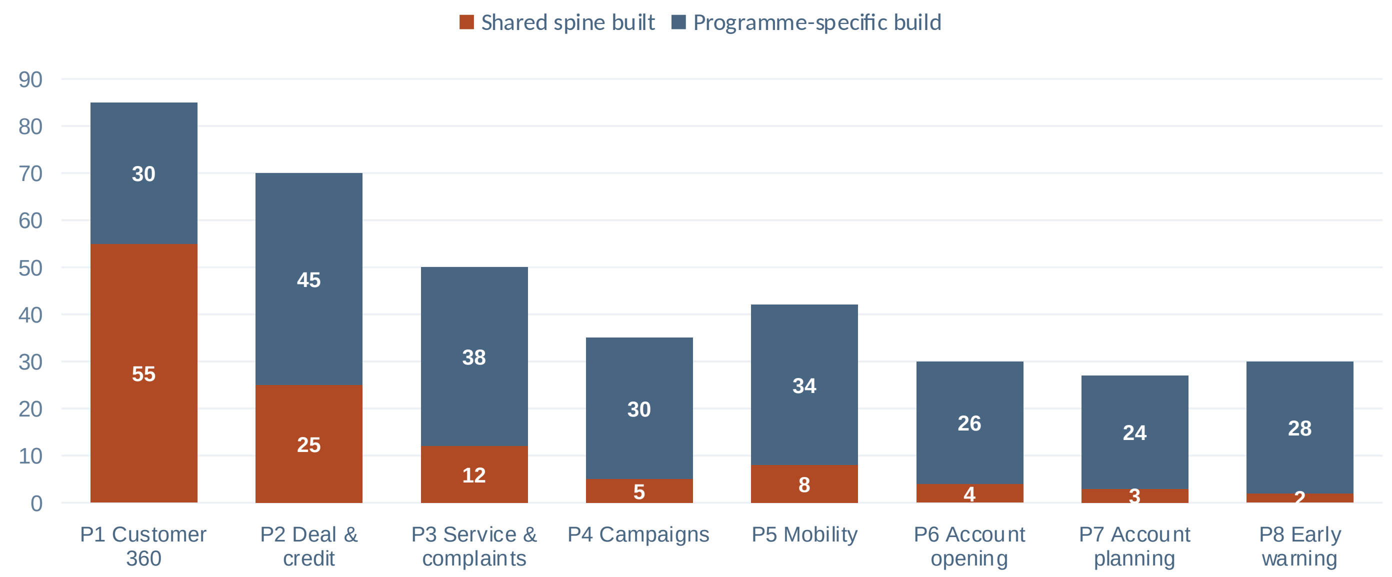

Illustrative effort index per programme — spine investment concentrated in the first two

85
effort index, programme one

30
effort index, programme eight

2
spines built, not eight

114
spine units invested total

What the curve does and does not show

- Programmes did not get simpler — programme eight was among the hardest in the portfolio.
- What fell was the share of each programme spent rebuilding things that already existed.

The argument this makes to a portfolio sponsor

Programme one carries roughly two-thirds of the total spine investment and shows the worst return of any programme in the portfolio. Judged individually — which is how programmes are usually judged — it is the weakest business case in the list.

That is the trap. A sponsor reviewing programme one on its own merits will cut the foundational scope, and programmes two through eight will each rebuild a portion of it at full price. Funding the sequence rather than the programme is the decision that unlocks the curve, and it can only be taken before programme one starts.

Five structural differences that change the data model and the approval path — none of which is a configuration setting

A sixth control function

Sharia governance is a distinct control function with its own review, its own authority and its own record. It is not a variant of compliance and cannot be modelled as one.

Approval model

Contract form is part of the decision

A facility is structured as a specific contract type — cost-plus sale, lease, partnership, agency. The structure must be approved as well as the credit, and the two approvals have different owners.

Data model

Asset-backed sequencing

Several structures require ownership or possession steps in a defined order. That creates timing dependencies and documentation states a conventional facility does not have.

Process model

Sector screening at origination

Prohibited activities are screened as an eligibility gate, not as a risk weighting. A screened-out counterparty is not a priced risk; it is ineligible.

Eligibility model

Profit, not interest

Pricing terminology, calculation and the treatment of late payment all differ. Late charges are not income, which changes both accounting and how collections behave.

Pricing model

Why this matters beyond Islamic banking

A platform designed for conventional lending can be configured to display different labels, but it cannot easily accommodate a sixth approver with independent authority, or an eligibility gate that removes a counterparty rather than pricing it, or a facility whose contract structure is itself an approved artefact.

Those are data-model and approval-model properties. The general lesson is the one worth carrying: when a market difference shows up as a new actor, a new gate or a new object, it is architecture. When it shows up as different wording on the same object, it is configuration. Confusing the two is how platforms get retrofitted badly.

Source: Author's delivery record, 2018–2020 engagement · author's analysis

Six shared assets carried the portfolio — and each one was built inside a programme that had its own reason to need it

Shared asset	What it provides	Built in	Reused by
Customer and group model	Entity, hierarchy, related parties, relationship-manager ownership	Programme 1 — Customer 360	**All seven**
Entitlement and access	Role-based access, data scoping, hierarchy-driven visibility	Programme 1 — Customer 360	**All seven**
Parallel approval engine	Multiple control functions deciding concurrently, each with its own record	Programme 2 — Deal and credit	**Three**
Case and SLA model	Queues, escalation ladders, service-level timers, dispositions	Programme 3 — Service	**Four**
Document and evidence store	One store, one retention policy, one audit surface across programmes	Programme 2 — Deal and credit	**Five**
Integration contracts	Governed interfaces to core, risk, document and analytics systems	Programmes 1 and 2	**All seven**

The rule that made this work

- No asset was built as a platform initiative. Each was built inside a programme that independently needed it.
- Each stayed funded by a business case and proven by a real use.
- A shared asset with no first customer is a design exercise.

The rule that protected it

- The second programme to need an asset extends it; it never forks it.
- That decision is made under delivery pressure, by someone accountable above one programme.
- Across eight programmes it was the single most repeated argument.

What it required of the team

- One team across three years — the continuity is not incidental to the result, it is most of it.
- A rotating team rebuilds rather than extends — extending needs knowing what exists.
- That is an argument for team stability that a resourcing model rarely makes.

Source: Author's delivery record, 2018–2020 engagement

What held across eight programmes, what did not, and what I would sequence differently on the next portfolio

What held

- Sequencing Customer 360 first despite it having the weakest standalone business case in the portfolio.
- Building the parallel-approval engine inside the programme that most needed it, rather than as a platform initiative.
- Keeping one team across three years — most of the compounding traces back to continuity rather than to design.
- Refusing the first fork. Each refusal was contested and each one held the model together.

What did not

- Control-function engagement was treated as a per-programme activity. Five functions were re-briefed eight times, which was avoidable and expensive in their time, not ours.
- Reporting was built per programme and only consolidated later — the one asset that should have been shared and was not.
- The compounding was never made visible to sponsors, so each programme was still justified as though it were the first.

What I would change

- Establish a standing control-function forum in programme one, not a per-programme engagement cycle.
- Treat reporting as a shared asset from the start — it has the same reuse profile as the customer model and none of the visibility.
- Publish the effort curve to sponsors each year. Invisible compounding gets cut in the first budget round that comes under pressure.
- Name the spine owner formally, rather than relying on the delivery lead having the standing to refuse forks.

The pattern behind all three columns

Everything that held was a decision about order or ownership. Everything that did not was something treated as belonging to a programme when it actually belonged to the portfolio — control-function engagement, reporting, and the story told to sponsors. The failure mode is consistent and easy to name in advance: whatever is shared but invisible will be rebuilt or cut.

Source: Author's synthesis

Four judgements that carry to any multi-programme portfolio

01	**Own the sequence, or nobody does**	In a multi-sponsor portfolio every sponsor optimises their own programme and the platform fragments quietly. The order is the highest-leverage decision available and it exists only once, before the first programme starts.
02	**Build shared assets inside a programme that needs them**	A shared asset with no first customer is a design exercise. Every reusable component should be funded by a business case and proven by a real use before anything else is allowed to depend on it.
03	**Distinguish a new actor from a new label**	When a market or regulatory difference introduces a new approver, a new gate or a new object, it is architecture. When it changes wording on an existing object, it is configuration. Getting that wrong is how platforms are retrofitted badly.
04	**Make the compounding visible**	An effort curve nobody sees is a saving nobody funds. Publishing reuse to sponsors each year is what prevents the foundational scope being cut in the first pressured budget round.

Eight programmes on one platform is an outcome. Eight programmes where the eighth cost a third of the first is a sequencing decision — taken once, at the start, by someone whose accountability sat above any single programme.