

Abu Dhabi Islamic Bank

Early warning without owning the model — governing a decision you did not calculate

A risk engine computes the signal. The relationship platform governs what a human does about it. Keeping those two things apart is the whole design — and collapsing them is the most common and most expensive mistake.



INTEGRATION SCOPE · WHOLESALE AND BUSINESS BANKING PORTFOLIOS

Risk signals surfaced in the relationship view · associated to the counterparty and group · human review, justification and approval before closure

Aditya Singh · Project Lead and Product Manager · early-warning integration and alert governance

Vendor, platform and system names removed. The integration scope is the client's; the separation principle, lifecycle design and noise analysis are mine.

The platform that shows a risk signal must not be the platform that computes it — and must be the only place the human response is recorded

Situation

A specialist risk engine produces early-warning signals on wholesale and business counterparties: deterioration indicators, behavioural triggers, exposure movements and external events.

Relationship managers work in a different system, where the customer, the group hierarchy and the relationship history live.

Complication

Signals in a risk system that relationship managers never open are not early warnings. They are a report.

The obvious fix — surface them in the relationship platform — creates a subtler problem. Once the signal is visible there, pressure builds to adjust it there: suppress a noisy trigger, tune a threshold for one portfolio, add a local override. Each is reasonable in isolation and each moves risk logic into a system with no model governance around it.

The separation

One system computes and owns the model. The other governs the human response and owns the record of it. Neither does the other's job, and the boundary is defended explicitly rather than assumed.

The relationship platform may filter what it displays and how it prioritises attention. It may never change what the signal says.

THE NUMBERS THAT MATTER

1

system computes

1

system governs response

0

thresholds in the CRM

3

human inputs required

4

lifecycle states

1

record per alert

Why the boundary erodes: every request to move logic across it is individually reasonable and locally cheap. The cost is cumulative and shows up years later, when nobody can say which system decided what, and model governance covers only half the logic that is actually running.

The integration scope was the client's; the separation principle, the lifecycle design and the noise analysis are mine

PROGRAMME SCOPE Client context The risk engine, the signal types and the governance obligations as defined by the bank. - External risk engine in place - Wholesale and business portfolios - Counterparty and group association - Service-based integration - Human review obligation - Audit and retention requirements	DELIVERY RECORD My ownership The separation principle, alert lifecycle, human-in-the-loop design and the delivery to production. - Separation-of-concerns design - Four-state alert lifecycle - Review, justification and approval model - Group and counterparty association - Noise and actionability analysis - Requirement through production	ILLUSTRATIVE Modelled Alert volume, disposition mix and noise figures illustrate the actionability argument. Directional, not measured. - Alert volume and disposition mix - Actionability decomposition - Threshold sensitivity illustration - Replace with portfolio actuals - Shape is the argument - Method transfers unchanged	OUT OF SCOPE Not claimed Excluded deliberately. An integration and governance pattern, not a risk-methodology or portfolio disclosure. - The bank's risk model or thresholds - Any counterparty or exposure data - Any credit or watchlist decision - Model performance or validation - Any confidential client document - Measured portfolio outcomes
---	---	---	--

THE PRINCIPLE

Two systems, one boundary — and everything that crosses it in the wrong direction is a governance failure disguised as a feature request

THE RISK ENGINE

Computes the signal

- Owns the model, its inputs, its thresholds and its calibration
- Owns model validation, performance monitoring and model risk
- Owns the definition of what constitutes a signal at all
- Emits an alert with a subject, a type, a severity and a basis
- Never records what anyone did about it



THE RELATIONSHIP PLATFORM

Governs the response

- Associates the alert to the counterparty, the group and the relationship manager
- Presents it in the context of everything else known about that relationship
- Requires a human review, a stated basis and an approval before closure
- Owns the complete record of what was decided, by whom, and why
- Never recalculates, re-scores or overrides the signal

The requests that erode the boundary, and the answer to each


“Can we suppress this trigger for one portfolio?” — No, but the platform can deprioritise how it is presented, which achieves the operational aim without changing the signal.

“Can we adjust the threshold here rather than raise a change with the risk team?” — No. A threshold changed outside model governance is model risk with no owner, and it will not appear in any validation.

“Can the relationship manager mark this as a false positive and have it stop firing?” — The manager can record that judgement, and it should feed back to the risk team as evidence. It must not silently alter future behaviour of the model.

— Every one of these is reasonable, cheap and locally correct. Refusing them consistently is the entire job.

Four states from arrival to closure — and closure requires three things a human has to supply, none of which the risk engine can produce



The three human inputs

- A judgement — is this material for this counterparty, in this context.
- A stated basis — the reasoning, which is the part an auditor and a successor both read.
- An action — what happens next, with an owner and a date.


Why closure needs a second person

- An alert closed by the relationship manager who owns the relationship is a self-assessment.
- Approval by a separate authority is what makes the closure a control rather than an administrative act.
- It also creates the sampling population an internal audit function actually needs.

What the design forbids

- Bulk closure. If a hundred alerts can be closed in one action, none of them was reviewed.
- Closure without a stated basis, however obvious the case appears.
- Auto-expiry. An alert that ages out silently is a control that reports itself as working.

An alert nobody acts on is worse than no alert — it consumes attention and it reports the control as operating



Illustrative disposition mix — the shape of the problem, not measured portfolio data

The feedback loop that keeps the boundary intact

Every disposition recorded in the relationship platform is evidence about the model — but it is evidence, not a control input. It goes back to the risk team as structured feedback on a defined cadence, and the risk team decides whether to recalibrate.

That loop is what makes the separation sustainable rather than merely principled. Without it, relationship managers watch the same unhelpful alert fire for two years, conclude the system is broken, and start pressing for local suppression — and the boundary erodes for entirely understandable reasons.

Reading the mix honestly

- A third of alerts are explained by relationship context the risk engine cannot see — a large exposure movement that is a known seasonal pattern, for example.
- Duplicates and data-quality alerts are not risk signals at all; they are integration and master-data problems arriving disguised as risk.
- Only about a third genuinely change what anyone does — and that is not a bad ratio for an early-warning system.

What to do with each category

- Context-explained alerts: feed the disposition back to the risk team as calibration evidence — never suppress locally.
- Duplicates: fix in the association logic, which is the platform's job and not a model

Six design decisions at the interface — each one either protects the boundary or quietly moves logic across it

Decision	What was chosen	What it prevents
Alert payload	Subject, type, severity, basis and timestamp — no raw model inputs or intermediate scores	The platform accumulating enough data to be tempted into recalculating
Association logic	Counterparty and group resolution owned by the platform, using its own master hierarchy	Two systems disagreeing about which entities form a group
Severity handling	Displayed and used for ordering; never recomputed, never locally overridden	Threshold drift outside model governance
Suppression	Presentation-layer only, reversible, logged, and visible as suppressed rather than absent	Alerts disappearing with no record that they ever fired
Disposition feedback	Structured, scheduled, one-way to the risk team as evidence — not a control input	Front-line judgement silently altering model behaviour
Record ownership	The response record lives with the platform; the signal record stays with the engine	An audit trail split across two systems with no join key

The one that matters most

Association logic. If the risk engine resolves groups its own way and the relationship platform resolves them another, the same alert attaches to different entities in the two systems — and every downstream number, from exposure roll-up to alert counts by relationship manager, disagrees.

Source: Author's delivery record, early morning integration programmes, author's design

That disagreement is discovered late, usually by someone reconciling two reports, and it is expensive to unpick because both systems are already trusted. Fixing group resolution ownership

Four rules for governing a decision you did not calculate

- | | | |
|----|---|--|
| 01 | Separate computation from response, explicitly | One system owns the model and its governance; the other owns the human response and its record. Write the boundary down, because it will be tested by reasonable requests within the first quarter. |
| 02 | Presentation may change, signal may not | Deprioritising, ordering and filtering are legitimate platform behaviours. Adjusting a threshold, suppressing a trigger or recomputing a score are not — they are model changes made outside model governance. |
| 03 | Close with a basis and a second person | A judgement, a stated reason and an approval by someone other than the relationship owner. Bulk closure and silent expiry turn an early-warning control into a control that reports itself as working. |
| 04 | Build the feedback loop or lose the boundary | Dispositions are evidence for the model owner, delivered on a cadence. Without that route, the people living with unhelpful alerts will eventually get local suppression, and they will be right to ask. |

The hard part was never the integration. It was holding a line that erodes one reasonable request at a time — and building the feedback loop that makes holding it sustainable rather than merely principled.