

Arab Bank

From dashboards to decisions — why a hundred reports is a symptom, not an achievement

Corporate profitability and portfolio analytics across four dashboard families. The delivery question was never how to build them; it was which decision each one changes, and what to do about the ones that change nothing.

4

dashboard families

100+

reports in the estate

22+

executive dashboards

1

question asked of every one

3

measures that change behaviour

50+

product categories to allocate across

DEPLOYMENT CONTEXT · PUBLICLY REPORTED

Client and portfolio profitability · transaction banking contribution · global group consolidation · risk-weighted return on assets

Aditya Singh · Senior Manager — profitability and portfolio analytics · 2023

A dashboard that does not change a decision is a cost — and most reporting estates grow because nobody is allowed to say so

Situation

A corporate bank with more than a hundred reports and over twenty executive dashboards across profitability, transaction banking, group consolidation and risk-weighted return.

Each was requested by someone, built to specification, and delivered.

Complication

Reporting estates grow monotonically. A report is cheap to request, expensive to maintain, and effectively impossible to retire — because retiring it requires someone to assert that nobody needs it, and nobody wants to be that person.

The result is an estate where genuine decision-support sits alongside reports that are opened once a quarter out of habit, and no one can tell which is which.

The method used

Ask one question of every dashboard: which decision does this change, and who takes it? A dashboard that cannot name a decision and an owner is a report, and reports belong somewhere cheaper.

Applied to the profitability build, that question reduced the specification substantially — and it changed which measures were built first.

THE NUMBERS THAT MATTER

1

question, asked of every view

4

families that survived it

3

measures that change behaviour

2

allocation choices that decide everything

1

definition set, agreed once

0

dashboards without a named owner

The judgement this argues: in analytics the hard problem is almost never the calculation. It is agreeing what a client is, which costs to allocate to them, and who is accountable for acting on the answer. Those three are political questions, and they are settled before the first query is written or they are settled repeatedly afterwards.

Estate scale is publicly reported; the design method and the allocation analysis are mine — no client or profitability data appears anywhere

PUBLICLY REPORTED

Deployment context

Reporting-estate scale as publicly reported for the deployment. Retained with attribution.

- 100+ reports in the estate
- 22+ executive dashboards
- Four dashboard families
- 50+ product categories
- Six-country corporate footprint
- Multi-currency operation

DELIVERY RECORD

My ownership

The profitability and portfolio analytics programme, including the decision-first specification method.

- Decision-first specification method
- Four dashboard families delivered
- Cost-allocation design
- Definition set and governance
- Risk-weighted return view
- Requirement through production

ILLUSTRATIVE

Modelled

Estate composition and allocation examples illustrate the method. No real figures are used.

- Estate composition analysis
- Allocation worked example
- Adoption pattern by view type
- Replace with bank actuals
- Shape is the argument
- Method transfers unchanged

OUT OF SCOPE

Not claimed

Excluded absolutely. No client, exposure, revenue or profitability data appears in this deck in any form.

- Any client or counterparty data
- Any revenue or profitability figure
- Any exposure or portfolio position
- The bank's cost-allocation policy
- Any confidential report or dashboard
- Measured business outcomes

THE FOUR FAMILIES

Each family answers a different question for a different decision-maker — and each one survived the decision test for a different reason


Client profitability Which relationships earn their capital? WHO DECIDES Relationship manager, segment head WHAT CHANGES Coverage intensity, pricing at renewal, whether to keep pursuing a relationship at current terms	Transaction banking Where does non-credit income actually come from? WHO DECIDES Product head, transaction banking WHAT CHANGES Product investment, cross-sell targeting, fee schedule changes	Global group What is our consolidated position on this group? WHO DECIDES Credit committee, group head WHAT CHANGES Group limit setting, concentration appetite, coordinated coverage across markets	Risk-weighted return Are we being paid for the capital we are using? WHO DECIDES Chief risk officer, finance WHAT CHANGES Capital allocation between segments, portfolio rebalancing, risk-adjusted pricing floors
---	--	--	--

Why naming the decision-maker matters as much as naming the decision

A dashboard with a named decision and no named owner drifts into being informational within two quarters. Nobody is accountable for acting on it, so nobody is accountable for whether it is right — and an unaudited dashboard degrades quietly as definitions and source systems change around it.

Requiring both, before build, does two things. It reduces the specification, because several requested views cannot name either. And it creates a maintenance owner by construction: the person who takes the decision has a direct interest in the view remaining correct.

Reporting estates only grow — a report is cheap to request, expensive to maintain, and socially impossible to retire



The reporting estate as publicly reported — reports, dashboards and the product set they cover

The intervention that works, and the one that does not

What does not work: an estate rationalisation exercise. It requires proving a negative about every report, it makes enemies in every function, and the estate regrows within a year because nothing changed about how requests are made.

What does work: change the entry condition. Every new view must name a decision and an owner before it is built, and every view carries a review date at which its owner reconfirms it or it lapses. Lapsing by default inverts the social problem — nobody has to argue for retirement, someone has to argue for retention. Applied at the point of request, the estate stops growing without anyone having to fight about the existing hundred.

Why the estate grows the way it does

- Requesting a report costs the requester nothing and signals diligence. Declining one costs the analytics team credibility.
- Each report is individually defensible. The estate is only indefensible in aggregate, and nobody owns the aggregate.
- Retiring a report requires asserting that nobody needs it — a claim that is unfalsifiable and personally risky.
- A hundred reports over four dashboard families is not four families' worth of questions. It is the same questions asked in twenty-five variations each.

Client profitability is an allocation argument wearing a calculation costume — two choices decide the answer and both are political

CHOICE ONE

What is a client?

- The legal entity, or the group it belongs to? A subsidiary can be unprofitable while its group is highly profitable.
- If the group, then a relationship manager covering one loss-making subsidiary appears to be destroying value.
- If the entity, then nobody is accountable for the group relationship and coordinated pricing becomes impossible.
- Most banks need both views. Few agree in advance which one the word profitability refers to by default.

CHOICE TWO

Which costs land on them?

- Direct costs are uncontroversial. Everything after that is a negotiation.
- Funding cost depends on a transfer-pricing method that finance owns and the business disputes.
- Capital cost depends on the risk weighting, which makes the same client look different under different regimes.
- Overhead allocation is where the argument becomes explicit — any basis chosen advantages some segments over others.

How this was handled, and why it is a delivery decision rather than a finance one

Both choices were settled before the first query was written, by finance and the business jointly, and written down as a definition set with a named owner. That sounds obvious and is routinely skipped, because it is slow and contentious and the build can technically start without it.

Starting without it has a specific failure mode. Analysts make reasonable local assumptions to keep moving; each dashboard embeds slightly different ones; and six months later three views of the same client disagree. At that point the disagreement is discovered by a relationship manager or a committee, in public, and the whole analytics layer loses credibility — not because any calculation is wrong, but because the definitions were never agreed.

The delivery lead's job here is not to resolve the allocation argument. It is to refuse to build until someone else has, and to make that refusal early enough that it costs a week rather than a programme.

Three measures moved decisions; the rest informed them — and the difference is whether the measure has a consequence attached

Measure	Why it changes behaviour	The consequence attached	Effect
Risk-weighted return	Makes the capital cost of a relationship visible against its income	Pricing floors and capital allocation between segments	Strong
Share of wallet	Reframes a large client as an underpenetrated one, which is a different conversation	Coverage intensity and cross-sell targets	Strong
Non-credit income share	Exposes relationships that look profitable only because of a single lending margin	Product investment and fee schedule review	Strong
Absolute revenue	Familiar and widely used, but rewards size rather than efficiency	None directly — informs, rarely decides	Weak
Client count by segment	Describes the portfolio without indicating what to do about it	None	None
Activity volumes	Measures effort rather than outcome; useful diagnostically, misleading as a target	None — and harmful if targeted	Negative

The test	The one to watch	What this implies for sequencing
• A measure changes behaviour when someone's decision, target or compensation depends on it. • Without a consequence even an excellent measure is read once. • That is not a failure of the dashboard — it is a failure to connect it to anything.	• Activity volumes are the classic trap: easy to measure, satisfying to display, and actively harmful once targeted. • Targeting calls made produces calls made, not revenue — and degrades the data. • Build it as a diagnostic, never as a target.	• Build the three measures with consequences first, hard as they are. • Descriptive views are easier and are what stakeholders ask for first. • Delivering the easy ones first frames analytics as reporting.

Four rules for building analytics people actually use

- | | | |
|----|---|---|
| 01 | Ask which decision changes, and who takes it | Both, before anything is built. A view that cannot name a decision is a report; a view that cannot name an owner has no maintenance sponsor and will degrade quietly as the definitions around it move. |
| 02 | Settle the definitions before the first query | What a client is and which costs land on it are political questions, not technical ones. Building before they are agreed guarantees that three views will disagree in public six months later. |
| 03 | Attach a consequence or expect no behaviour change | Measures move behaviour when a decision, target or compensation depends on them. Without that, even an excellent measure is read once. And never target an activity volume. |
| 04 | Change the entry condition, not the estate | Rationalisation exercises fail and the estate regrows. Requiring a named decision and owner at request, with lapse-by-default review dates, stops the growth without a fight over the existing hundred. |

The calculation was never the hard part. Agreeing what a client is, which costs belong to it, and who is accountable for acting on the answer — those are the three decisions that determine whether an analytics layer is used or merely delivered.