

Maybank

What a technical panel actually asks — the agenda is the scoring sheet

A full-day technical and security evaluation is not a product demonstration. The session list is the scoring sheet, each session tests something different from what its title suggests, and most responses fail in the same five places.

6

evaluation sessions

1

full day, two parallel tracks

5

questions that decide the day

5

recurring failure modes

0

marks for assertion

1

trap almost everyone falls into

EVALUATION FORMAT · AS SET BY THE BANK

Solution architecture · scalability · integration · data migration · IT security · customisation and extensibility

Aditya Singh · Analysis authored to prepare and structure the technical evaluation response · 2026

Each session tests something other than its title — and panels award marks for evidence, never for assertion

Situation

A full-day technical and IT-security evaluation with six sessions: solution architecture, scalability, integration and data exchange, data migration, security and compliance, and customisation and extensibility.

Attended by the bank's architects, security function and platform engineers, running two parallel tracks with a vendor helpdesk alongside.

Complication

Most suppliers prepare six presentations matching six titles. That is the wrong preparation.

The architecture session tests whether the architecture is real or a diagram. The scalability session tests whether anyone has measured anything. The migration session tests honesty. And a single stale version number in an appendix can undo a strong morning, because it tells the panel how current the platform actually is.

The method

Prepare against what each session tests rather than what it is called. Bring evidence rather than assertions — measured numbers, real diagrams, actual failure stories.

And audit every artefact for currency before the day, because the panel reads the appendices even when the room does not.

THE NUMBERS THAT MATTER

6

sessions to prepare

6

different things tested

5

questions that decide it

5

recurring failure modes

1

currency audit, always

0

marks for assertion

The judgement this argues: a technical panel is not deciding whether the platform can do the job — the shortlist already established that. It is deciding whether the people in the room have run this before, and every question is designed to find out.

The session structure is the client's; the analysis of what each session tests and the preparation method are mine

EVALUATION FORMAT Client context The agenda, session list, timing and attendee composition as set by the bank for the evaluation day. - Six sessions across a full day - Two parallel tracks - Vendor helpdesk alongside - Architecture and security focus - Clarification windows scheduled - Attendee mix by function	ANALYSIS & METHOD My contribution What each session actually tests, the five deciding questions, the failure modes and the preparation discipline. - Session-to-test mapping - Five deciding questions - Five recurring failure modes - Evidence-over-assertion standard - Currency audit method - Room and role allocation	PATTERN Observed Failure modes drawn from technical evaluations across multiple enterprise pursuits, not from a single event. - Recurring across pursuits - Not attributed to any one panel - No client-specific outcomes - Generalised deliberately - Method transfers unchanged - Not a scoring reconstruction	OUT OF SCOPE Not claimed Excluded deliberately. This is a preparation method, not a record of any evaluation or its result. - Any panel's actual questions - Any evaluation score or outcome - Whether the pursuit was won - Any confidential client material - Any competitor's response - Any product's certified capability
---	---	--	--

Six sessions, six different tests — and in every case the title describes the topic while the session measures something else entirely

Session	What it actually tests	The tell	Weight
Solution architecture	Whether the architecture is real or a diagram drawn for the bid	Can they answer a question the diagram does not cover	Decisive
Scalability	Whether anyone has measured anything, or is quoting a datasheet	Do the numbers come with a test method attached	Decisive
Integration and data exchange	Whether integration is a product capability or a services engagement	How many interfaces exist versus how many are built per client	High
Data migration	Honesty — this is the session where over-claiming is cheapest to detect	Do they volunteer what should not be migrated	High
Security and compliance	Whether security is designed in or certified around	Is the answer a control or a certificate	Decisive
Customisation and extensibility	Whether configurability is real or a screen in front of code	Ask for a change the demonstration script did not anticipate	High

Why the titles mislead	What the panel is really deciding	How that changes preparation
- Titles are written by procurement to structure a day, not to structure a test. - The questions come from architects, testing something narrower and more specific — so preparing to the title yields six competent presentations that answer nothing asked.	- Not whether the platform can do it — shortlisting already settled that. - Whether the people in the room have run this before, at this scale, and know where it hurts. - Every question probes operating experience, not product knowledge.	- Bring the person who ran it, not only the person who can present it. - Prepare evidence for each test, not slides for each title. - Rehearse the questions the diagram does not answer, because those are the ones that get asked.

THE FIVE QUESTIONS

Five questions decide the day — none is about features, and all five have a good answer that most teams have not prepared

01

What breaks first, and at what number?

Every platform has a first bottleneck. A team that names theirs — with the number and the mitigation — is telling the panel it has run this at scale. A team that says nothing breaks is telling them it has not.

The number, the symptom, and what you do about it

02

Show me an interface you did not build for this bid

Tests whether integration is a product capability or a services engagement dressed as one. An adapter inventory with versions and clients behind it answers this; a diagram does not.

A real inventory, with what is prebuilt versus per-client

03

What would you refuse to migrate?

The honesty probe. A supplier who wants to migrate everything has not done a large migration. The right answer names the categories that should stay in the legacy system and explains why.

A decision tree, and the categories you would leave

04

Walk me through your last serious incident

Tests operating maturity more directly than any architecture question. Detection time, decision path, communication, root cause and what changed afterwards.

One real incident, told with the uncomfortable parts intact

05

Make a change we have not scripted

The configurability test. Ask for something adjacent to the demonstration and watch whether it is configuration or a conversation about a release.

Do it live, or say plainly what it would take

Five ways strong platforms lose technical evaluation days — four are preparation failures and the fifth is a hygiene failure

01 The diagram with no author in the room Architecture is presented by someone who did not design it. The first question outside the diagram exposes it, and everything afterwards is discounted. THE FIX Bring the designer. If they cannot attend, do not present their diagram as your own.	02 Datasheet numbers with no method Concurrency and throughput figures quoted without the test conditions that produced them. An engineering panel treats an unattributed number as no number. THE FIX Quote fewer numbers, each with its load profile, environment and date.	03 Certificates instead of controls Security answered with a list of accreditations rather than with the control that stops the specific attack asked about. THE FIX Lead with the control, cite the certificate as corroboration afterwards.	04 Migrating everything An offer to migrate the full history, unprompted. Reads as inexperience to anyone who has run a large migration and had to defend the data quality afterwards. THE FIX Bring a migration decision tree and the categories you would refuse.	05 Stale versions in the appendix Runtime and orchestration versions several years past end of life, sitting two slides from a current application framework. The panel reads the appendix. THE FIX Audit every version number in every artefact the week before. Every time.
---	--	--	--	--

Why the fifth is worth its own discipline

It is the only one that costs nothing to fix and the only one that undermines everything else. A panel that finds a runtime version five years past end of life in an appendix will re-read the whole response asking what else is out of date — and the architecture session, however strong, is now being assessed by someone with a specific doubt.

An artefact currency audit the week before every evaluation takes an afternoon. It is the highest-return preparation activity available and the one most reliably skipped.

For each session, one artefact beats a presentation — and the artefact is usually something that already exists internally

Session	The artefact to bring	Where it comes from	Effort
Solution architecture	A deployment diagram from a live client, redacted — not a reference architecture	Existing production documentation	Low
Scalability	A load-test report with profile, environment, date and the bottleneck found	Last performance test cycle	Low
Integration	The adapter inventory: what is prebuilt, what is configured, what is built per client	Product catalogue	Low
Data migration	A migration decision tree and a real staged plan with reconciliation points	A prior migration's runbook	Medium
Security	The control matrix mapped to attack classes, with certificates as corroboration	Existing control documentation	Medium
Extensibility	A recorded configuration change, timed end to end including approval	One afternoon in a sandbox	Medium

Why artefacts outperform slides	The one worth building fresh	What to leave behind
• It was produced for a purpose other than winning this deal — which is why a panel trusts it. • It survives the room and is read again during scoring. • And it is much harder to fake.	• The timed configuration change is the only row needing new work, and it is the highest-value one. • A recording with a visible clock, covering approval as well as configuration, answers the extensibility question completely. • One afternoon in a sandbox, reusable across every subsequent pursuit.	• Reference architectures that no client actually runs. • Certification lists presented as security answers. • Any number without a method attached to it.

Source: Author's analysis of technical evaluation days

Six sessions, two tracks, one room to manage — and the most common structural error is sending the same three people to everything

The architect who built it Owns architecture, integration and extensibility. Must be someone who has made trade-offs on a live deployment, not a presenter briefed on the diagram.	The engineer who broke it Owns scalability and incidents. The credibility of the whole day rests on someone being able to describe a real failure without flinching.	The security owner Owns the security and compliance track. Answers with controls first and certificates second, and can go deeper than the deck on any single control.	The migration lead Owns data migration. The one person in the room whose job is to say what should not be moved.	The deal lead, silent Present, tracking questions and commitments, intervening only on scope and commitment — never answering a technical question on someone else's behalf.
---	---	---	---	---

Handling a question you cannot answer

- Say so, name who can, and commit to a written answer with a date — then meet it.
- Panels expect gaps. They do not expect a confident wrong answer, and they remember one.
- A tracked question log, shared back the same evening, is worth more than an improvised answer.

Handling the parallel track

- Two tracks means two panels forming independent impressions, and they will compare notes.
- An inconsistency between tracks reads as either disorganisation or evasion.
- Brief both teams on the same answers to the five deciding questions before the day starts.

The last thirty minutes

- The clarification window is not a formality — it is the last chance to correct a misread.
- Come with the two things you most want restated, and say them.
- Close by confirming the open-question list and when each answer lands.

Four rules for any technical evaluation day

- | | | |
|----|---|--|
| 01 | Prepare to the test, not to the title | Session titles are written to structure a day; the questions come from architects testing something narrower. Map each session to what it actually measures before preparing a single slide. |
| 02 | Bring artefacts, not assertions | A redacted production diagram, a dated load-test report and a timed configuration recording outperform any presentation — largely because they were produced for a purpose other than winning this deal. |
| 03 | Send the people who did it | The panel is assessing operating experience, not product knowledge. Someone who can describe a real incident with the uncomfortable parts intact is worth more than a polished presenter. |
| 04 | Audit currency before every evaluation | One stale version number in an appendix invites a panel to re-read the whole response looking for what else is out of date. An afternoon's audit is the highest-return preparation there is. |

A technical panel is not deciding whether the platform can do the job — the shortlist settled that. It is deciding whether the people in the room have done it before, and every question is built to find out.