Notes from 38 Programmes

What actually predicts whether a banking transformation works

Thirty-eight programmes delivered across sixteen institutions, plus solution and analysis work at nine more. Four regions, eight years. Almost none of what separated the programmes that worked from the ones that struggled was technology — and the pattern is consistent enough to be worth stating plainly.

38	16	9	4	8	4
programmes delivered	institutions delivered for	more, solution and analysis	regions	years	decisions that decide it

PORTFOLIO SCOPE

Delivery at sixteen institutions · solution, presales and analysis at nine more · retail, corporate, wholesale, Islamic, insurance and community banking

Aditya Singh · Senior Manager — Global Operations and Business Technology · BFSI delivery 2018–2026

Client and vendor names removed where the argument does not require them. Every finding is drawn from the twenty-nine case decks in this portfolio.

Across thirty-eight programmes the same four decisions decided the outcome — and all four are taken in the first quarter, before anything is built

Situation

Thirty-eight programmes delivered across sixteen institutions: a Gulf Islamic bank, India's largest private bank, a regional corporate bank across six countries, an Indian life insurer, and twelve American credit unions and community banks.

Alongside that, solution, presales and analysis work at nine more — Indian public sector, Southeast Asian universal and digital-first banks, and an insurance distribution pursuit.

Complication

The programmes that struggled did not struggle on capability. In almost every case the platform could do what was asked, the team was competent, and the requirements were reasonable.

They struggled because a decision taken early and cheaply — usually without anyone noticing it was a decision — made everything afterwards more expensive. And the cost only became visible in year two or three, by which point it was structural.

The finding

Four decisions recur across every programme in the portfolio regardless of region, segment or scale: what the shared object is, in what order things are built, who is permitted to say no, and what gets measured before anything changes.

Get those four right and ordinary teams deliver well. Get them wrong and strong teams spend three years compensating.

THE NUMBERS THAT MATTER

4	**1**	**0**	**5**	**5**	**8**
decisions that recur	quarter in which all four are taken	that are technology choices	findings behind them	failure modes that repeat	years, one pattern

Why this is worth stating as a pattern rather than as anecdotes: each of the four is individually obvious. What is not obvious is that they are the same four every time, that they are all taken before the first sprint, and that none of them is the decision anyone convenes a steering committee to take.

This deck argues from the other twenty-nine — every finding traces to a specific programme, and the boundaries set in those decks apply here unchanged

DELIVERY RECORD

Portfolio context

The engagements the findings are drawn from, as set out in the individual case decks.

- 38 programmes delivered, 2018–2026
- 16 institutions delivered for
- 9 further institutions analysed
- Retail, corporate, wholesale, Islamic
- Life insurance and community banking
- Delivery and analysis stated separately

SYNTHESIS

My contribution

The pattern itself — the four decisions, the five findings and the five failure modes — and the argument that they are the same everywhere.

- Four-decision framework
- Five recurring findings
- Five recurring failure modes
- Cross-region pattern argument
- Applied prospectively since
- Traceable to named case decks

FROM THE CASE DECKS

Inherited

Figures cited here carry the same status they carry in the deck they come from — public, reported, modelled or reconstructed.

- Public figures stay attributed
- Modelled figures stay labelled
- Reconstructed detail stays flagged
- Nothing is upgraded by aggregation
- See the source deck for each
- Claims register available

OUT OF SCOPE

Not claimed

Excluded deliberately, as in every deck in this portfolio.

- Sole responsibility for any outcome
- Institution financial results
- Any confidential client material
- Independently audited findings
- Statistical generality beyond BFSI
- Commercial terms of any engagement

Source: Author's synthesis across twenty-nine case decks in this portfolio

Sixteen institutions delivered for and nine more analysed — deliberately different enough that a pattern holding across all of them is worth something

Delivered — Middle East
2 institutions

- Gulf Islamic bank — eight programmes across wholesale and business banking
- Regional corporate bank — four programmes across six countries

Delivered — India
2 institutions

- Largest private bank — six digital journeys
- Life insurer — policy servicing across two programmes

Delivered — United States
12 institutions

- Credit unions and community banks, roughly $0.5bn to $3.6bn in assets
- One productised delivery platform across the portfolio

Solution and analysis
9 institutions

- Indian public sector banking at scale
- Insurance distribution pursuit
- Southeast Asian universal, digital-first and reference deployments

Delivery versus analysis

- The sixteen are engagements I delivered on. The nine are solution, presales and analysis work — real, and a different class of claim.
- A pattern holding in a $3.6bn credit union and a bank with 300,000 platform users is unlikely to be an artefact of scale.
- Findings are drawn from both, and each case deck states which it is.
- The claim is not statistical representativeness — only that the base is varied enough to be interesting.

What was genuinely different

- Regulatory posture, product structure, distribution model, scale, and what counts as a customer.
- Islamic banking adds a control function and an approved contract object. A cooperative changes what segmentation even means.
- Those differences are real and they are the subject of individual decks.

What was the same everywhere

- The four decisions. Not their answers — their existence, their timing, and the cost of getting them wrong.
- Also the failure modes, which repeated with almost no variation across regions.
- That consistency is the finding worth carrying.

Source: Author's synthesis across the portfolio

WHAT PREDICTED SUCCESS

Five findings that recurred across every region — and each traces to specific programmes rather than to a general principle

01	**Sequence beat capability, every time**	The programmes that compounded were the ones where someone owned the order. Foundational work first, even where it had the weakest standalone business case, because everything after it got cheaper.	*Gulf portfolio · multi-country corporate · US factory*
02	**Reuse survived only where someone could refuse a fork**	Shared components stayed shared because one person's accountability sat above any single delivery date. Where it did not, the platform quietly became several platforms within two years.	*Digital journeys · eight-programme portfolio · deployment factory*
03	**The binding constraint was rarely the one in the brief**	Briefs describe capability gaps. The actual constraint was usually assurance capacity, change-approval throughput, data completeness or a counterparty — and none of those appear in a requirements document.	*Enterprise scale · configurability · corporate spine*
04	**Controls designed in were a speed strategy**	In every large estate, assurance capacity was scarcer than engineering capacity. Journeys and programmes that arrived at review with controls already designed shipped sooner than those that arrived fast and unfinished.	*Digital journeys · parallel approvals · three regulators*
05	**Measurement before change made the result defensible**	Where a baseline was captured before release, the outcome could be attributed. Where it was not, a real improvement was indistinguishable from the quarter — and got argued about instead of banked.	*Journey portfolio · case visibility · analytics*

Source: Author's synthesis across the portfolio

WHAT PREDICTED STRUGGLE

Five failure modes repeated with almost no variation — all five are cheap to prevent and structural by the time they are visible

01	**Building for the hardest case, then applying it universally**	A control necessary for a minority becomes the process for everyone. It deters exactly the cases you want and adds nothing to the ones it was written for.	Deters the majority
02	**Conflating two states that needed separating**	Waiting on us and waiting on them. Direct and indirect exposure. Signal and response. Each conflation produces a confident number nobody can act on.	Measurement dies
03	**Deferring the irreversible decision**	Residency, tenancy, the object model, the hierarchy version. Each is cheap on day one and effectively impossible on day one thousand, and none of them demonstrates well.	No recovery later
04	**Leaving shared things invisible**	Reporting definitions, control-function engagement, the reuse curve itself. Whatever is shared but unowned gets rebuilt by each programme or cut in the first pressured budget round.	Rebuilt or cut
05	**Handing down a target instead of designing one in**	First-touch resolution, deflection, adoption, activity volume. A number given to a front line without the authority, data and tooling behind it produces the number and not the outcome.	Metric is gamed

Source: Author's synthesis across the portfolio

Every finding and every failure mode reduces to one of four questions — asked in the first quarter, answered by default if nobody asks them

01
What is the shared object?

The customer, the group, the case, the facility — the thing every programme will attach to.

WHY IT DECIDES THE OUTCOME

Get it wrong and every subsequent programme reconciles rather than builds. It is the only decision that is both invisible and load-bearing.

Group not country. Member and household. Boundary as the design unit.

02
In what order?

Which programme goes first, which regulator is engaged first, which market leads.

WHY IT DECIDES THE OUTCOME

Nobody owns the order in a multi-sponsor portfolio, so it defaults to whoever has the strongest business case — which is systematically the wrong answer.

Foundations first. Hardest regulator first. Taxonomy before interface.

03
Who can say no?

Who is permitted to refuse a fork, a threshold change, a private status vocabulary, on the day a delivery date is at risk.

WHY IT DECIDES THE OUTCOME

If that authority sits inside a single delivery it will lose every argument it has, because it is being asked to prioritise something invisible over something dated.

Named spine owner. Refuse the first fork. Gate, do not weight.

04
What is measured before it changes?

The baseline, captured before release, against a metric someone owns.

WHY IT DECIDES THE OUTCOME

Without it a real improvement is indistinguishable from the quarter, and the programme spends its review arguing rather than reporting.

Baseline per journey. Owner metric per objective. Leading before lagging.

None of the four is a technology choice. All four are taken in the first quarter, at almost no cost, and none of them gets a steering committee.

Source: Author's synthesis across the portfolio

What I would do in the first ninety days of any programme of this shape

01 **Name the shared object and write it down**

Before scope, before architecture. Whether it is the group, the member, the household or the case — every later argument about duplication, reconciliation and reporting traces back to whether this was settled once or assumed four times.

02 **Take ownership of the order**

In a multi-sponsor portfolio the sequence defaults to whoever has the strongest business case, which is systematically wrong because foundational work never does. Owning the order is the highest-leverage authority available and it exists only at the start.

03 **Establish who can refuse, above any delivery date**

A named owner of the shared spine whose accountability does not sit inside one programme. Without it, the first fork happens under pressure in month three and every subsequent one cites the precedent.

04 **Instrument before the first release**

A baseline per outcome, an owner per metric, leading indicators agreed before the pilot review. It costs a fortnight and it is the difference between a programme that can prove its result and one that has to argue for it.

Thirty-eight programmes delivered across sixteen institutions, nine more analysed, four regions, eight years — and the same four questions every time. None is about the platform, all four are answered in the first quarter, and the ones answered by default are the ones that cost three years.

Source: Author's synthesis across the portfolio